

# **Data Analysis**

# **Report**

**Click To Cart**

**A predictive modeling task**

**Name - Roshni Laha**

**Roll no. - 24IE10002**

# EDA : Visualization of data and correlation with target column

## Features in the data :

```
df.info()  
  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 1800 entries, 0 to 1799  
Data columns (total 8 columns):  
 #   Column            Non-Null Count  Dtype     
---  --     
 0   Time_on_site      1800 non-null   float64  
 1   Pages_viewed     1800 non-null   float64  
 2   Clicked_ad       1800 non-null   int64  
 3   Cart_value       1800 non-null   float64  
 4   Referral          1800 non-null   object  
 5   Browser_Refresh_Rate 1800 non-null   float64  
 6   Last_Ad_Seen     1800 non-null   object  
 7   Purchase          1800 non-null   int64  
dtypes: float64(4), int64(2), object(2)  
memory usage: 112.6+ KB
```

Fig1 : Showing column names, count, and Data type

The dataset has :

- o 4 columns of float type
- o 2 columns of int type (including target column ‘Purchase’)
- o 2 columns of object type

There are no null values in any column

```
df.isnull().sum()  
  
Time_on_site    0  
Pages_viewed    0  
Clicked_ad      0  
Cart_value      0  
Referral         0  
Browser_Refresh_Rate 0  
Last_Ad_Seen    0  
Purchase         0  
  
dtype: int64
```

Fig2 : Showing count of null values in each column

There are no duplicate rows

```
df.duplicated().sum()
```

```
np.int64(0)
```

Fig3 : Showing there are no duplicate rows

## Plots




Fig4 : Scatterplot showing linear relationship between Time on site and Pages Viewed

```
[1] sns.boxplot(x='Clicked_ad',y='Pages_viewed',data = df,hue='Purchase')
plt.xlabel('Clicked_ad')
plt.ylabel('Pages_viewed')
plt.show()
```




Fig5 : Box Plot between Pages viewed and Clicked ad

```
[15] sns.displot(data=df,x='Time_on_site',hue='Purchase',kind='kde')
```

```
[16] <seaborn.axisgrid.FacetGrid at 0x7f7e652f6890>
```




Fig6 : KDE plot of Time on site




Fig7 : KDE plot of cart value




Fig8 : KDE plot of Browser Refresh Rate

## Categorical columns

```
[31] df['Referral'].value_counts().plot(kind='pie', autopct='%.2f')
```

```
↳ <Axes: ylabel='count'>
```




Fig9 : Showing percentage of referrals

```
[117] df_refers
```

```
↳
```

```
Referral_count Referral_count_for_purchase_true Percentage_who_purchased
```

### Referral

Direct	347	107	30.84
Facebook	360	111	30.83
Google	741	205	27.67
Instagram	352	103	29.26

```
[138] df_refers['Percentage_who_purchased'].plot(kind='pie')
```

```
↳ <Axes: ylabel='Percentage_who_purchased'>
```




Fig10 : Showing percentage of people of each category who purchased

```
df['Last_Ad_Seen'].value_counts().plot(kind='pie', autopct='%.2f')  
<Axes: ylabel='count'>
```




Fig11 : Showing percentage of each category of last ad seen

```
[120] df_last_ads  
[120] Last_ad_seen  Last_ads_count  Last_ads_count_for_purchase_true  Percentage_who_purchased  
Last_Ad_Seen  
A 421 120 28.50  
B 452 138 30.53  
C 481 131 27.23  
D 446 137 30.72
```

```
df_last_ads['Percentage_who_purchased'].plot(kind='pie')  
<Axes: ylabel='Percentage_who_purchased'>
```




Fig12 : Showing percentage of each category who purchased

## Correlation with Target Column




Fig13 : Showing the count of target variable - Purchase

```
sns.boxplot(x='Purchase',y='Time_on_site',data = df)
plt.xlabel('Purchase')
plt.ylabel('Time_on_site')
plt.show()
```




Fig14 : Box Plot between Purchase and Time on site

```
sns.boxplot(x='Purchase',y='Pages_viewed',data = df)
plt.xlabel('Purchase')
plt.ylabel('Pages_viewed')
plt.show()
```




Fig15 : Box plot between Purchase and Pages viewed

```
sns.boxplot(x='Purchase',y='Browser_Refresh_Rate', data=df)
plt.xlabel('purchase')
plt.ylabel('Browser_Refresh_Rate')
plt.show()
```




Fig16 : Box Plot between Purchase and Browser Refresh Rate

## Key Observations :

- The dataset exhibits a class distribution of approximately 70% for Purchase = 0 and 30% for Purchase = 1. (**slightly imbalanced**)
- Time on site and pages viewed are linearly related (*Fig 4*).
- People who have purchased have more Time on site and more Pages viewed (*Fig 6 , Fig 14 , Fig 15*).
- Browser refresh rate** is a not so important feature for the target column as seen from the box plot (*Fig 8, Fig 16*).
- Google has the largest (41.17%) referral.
  - Nearly 30% of people from each Referral category have made a purchase (*Fig 10*).
- All categories of Last ad seen have nearly equal (25%) count.
  - Nearly 30% of people from each Last ad seen category have made a purchase (*Fig 12*).
- Column Transformations :
  - Browser\_Refresh\_Rate column is dropped** as it is quite irrelevant.
  - Referral and Last\_Ad\_Seen** (categorical columns) are converted to numerical by one hot encoding.
  - The first one hot encoded column is not dropped. Dummy variable trap is overlooked as we are working with tree-based algorithms (AdaBoost and XGBoost) and no matrix inversion operations are involved.

# Model Selection and Explanation

---

The models selected are **AdaBoost and Xg Boost.**

## **AdaBoost :**

- It is a stagewise additive model.
- It uses **Decision Stumps** as its base model.
  - Decision Stumps are weak learners (Decision Trees with max\_depth=1)

## **Overall Working of AdaBoost Classifier (Big Idea) :**

- This algorithm first trains a base classifier (a decision stump) and makes predictions on the training set.
- The algorithm then **increases the relative weight of misclassified training instances.**
- Then it trains a second classifier, using the updated weights, and makes predictions on the training set.
- Again increases the relative weight of misclassified training instances.
- and so on....
- After all the predictors are trained, the ensemble makes predictions by **adding the predictions of each predictor multiplied by the weight of the predictor.**
  - Each predictor has a weight depending on their overall accuracy on the weighted training set.

## Formulas :

Model :

$$h(x) = \sum_{i=1}^m (\alpha_i * h_i(x))$$

If  $h(x)$  is +ve then Purchase = 1 else Purchase = 0

- $m$  is the number of predictors.
- $\alpha$  is the weight of each predictor.

$\alpha$  (weight of each predictor) :

$$\alpha = 0.5 * \ln\left(\frac{1 - error}{error}\right)$$

Error :

error =  $\sum$  weight of  
misclassified training  
instances.

Weight of each training instance initially =  $\frac{1}{no. of rows}$

New weight :

For misclassified points	$New\_wt = current\_wt * e^{\alpha_m}$
For correctly classified points	$New\_wt = current\_wt * e^{-\alpha_m}$

$\Sigma weights must be = 1 so weights must be normalised.$

$$\text{Normalised weight} = \frac{weight}{\Sigma weight}$$

### Assumptions of the model :

- The dataset does not have missing values.
- The dataset has less outliers as adaboost has the tendency of overfitting.

### The model makes prediction using the formula :

$$h(x) = \sum_{i=1}^m (\alpha_i * h_i(x))$$

Where m is the number of predictors.

### This model is used because :

- Our dataset does not have any missing value.
- The Browser\_Refersh\_Rate feature has been dropped as it was irrelevant . It was also a noisy feature. So, the chances of overfitting are less.
- Moreover AdaBoost is an efficient boosting algorithm, which works great on small datasets like our customerBehaviour dataset.

### For this dataset the best hyperparameters of AdaBoost are :

- n\_estimators = 10
- learning\_rate = 0.5

**Accuracy = 79.56%**

**Precision score 0.72**

**Recall score 0.54**

**F1 score 0.62**

**AUC = 0.86**

### **Confusion matrix :**

	<b>0</b>	<b>1</b>
<b>0</b>	285 (True -ve)	29 (False +ve)
<b>1</b>	63 (False -ve)	73 (True +ve)

**The AdaBoost model from scratch gives an accuracy of 76.89% .**

## XG Boost :

Formulas :

Prediction :

$$\hat{y}_i = \sum_{k=1}^K f_k(x_i)$$

We add such  $f_k$  after each iteration so that objective function is minimised.

where  $f_k$  is a regression tree model.

Objective function :

$$\mathcal{L}^{(t)} = \sum_{i=1}^n L(y_i, \hat{y}_i^{(t-1)} + f_t(x_i)) + \Omega(f_t)$$

$L$  is the loss function.

- MSE (mean squared error) for regression problems
- Log\_loss for classification problems

$\hat{y}_i^{(t-1)}$  is the prediction of the  $i^{\text{th}}$  instance after  $(t - 1)^{\text{th}}$  iteration.

$\Omega(f_t)$  is a built in regularization parameter

$$\Omega(f) = \gamma T + \frac{1}{2} \lambda ||w||^2$$

- $T$  is the number of leaf nodes.
- $w$  is the weight/output of leaf nodes.
- $\gamma$  and  $\lambda$  are regularization parameters.

$f_t$  is the last added DT after  $t^{\text{th}}$  iteration.

We can approximate the loss function part of objective function, to make it differentiable (so that we can carry out optimizations), using Taylor series

about the point  $a = \hat{y}_i^{(t-1)}$ .

Objective function after approximation :

$$\mathcal{L}^{(t)} = \sum_{i=1}^n \left[ L(y_i, \hat{y}_i^{(t-1)}) + g_i f_t(x_i) + \frac{1}{2} h_i f_i^2(x_i) \right] + \Omega(f_t(x_i))$$

$g_i = \frac{\partial L(y_i, \hat{y}_i^{(t-1)})}{\partial \hat{y}_i^{(t-1)}} \text{ (gradient)}$	$h_i = \frac{\partial^2 L(y_i, \hat{y}_i^{(t-1)})}{\partial \hat{y}_i^{(t-1)2}} \text{ (hessian)}$
--------------------------------------------------------------------------------------------------	----------------------------------------------------------------------------------------------------

$L(y_i, \hat{y}_i^{(t-1)})$  is a constant , so removing it and expanding  $\Omega(f_i(x_i))$  :

Final simplified Objective function :

$$\Rightarrow \mathcal{L}^{(t)} = \sum_{j=1}^T \left[ \sum_{i \in I_j} g_i w_j + \frac{1}{2} \left( \sum_{i \in I_j} h_i + \lambda \right) w_j^2 \right] + \gamma T$$

For **objective function to be minimum**  $\frac{\partial L}{\partial w_j} = \mathbf{0}$  . This gives

$$w_j = - \frac{\sum_{i \in I_j} g_i}{\sum_{i \in I_j} h_i + \lambda}$$

Substituting this value in final objective function we get :

$$\mathcal{L}^{(t)} = \sum_{j=1}^T \left[ - \frac{1}{2} \frac{\left( \sum_{i \in I_j} g_i \right)^2}{\sum_{i \in I_j} h_i + \lambda} \right] + \gamma T$$

Also called similarity score

Loss reduction after a split is given by :

$$\text{Similarity score}_{\text{left child node}} + \text{Similarity score}_{\text{right child node}} - \text{Similarity score}_{\text{root node}}$$

### Assumptions of the model :

XGBoost is a scalable tree boosting algorithm.

- The loss function must be differentiable.
- There must be less outliers (to avoid overfitting).
- The dataset must not be very small.

### The model makes prediction using the formula :

$$\hat{y}_i = \sum_{k=1}^K f_k(x_i)$$

### This model is used because :

- XGBoost handles class imbalance (our model is slightly imbalanced).
- It provides high accuracy.
- It supports parallel processing and uses cache memory efficiently.

### For this dataset the best hyperparameters of XGBoost are :

- n\_estimators = 20
- learning\_rate = 0.1
- colsample\_bytree = 0.8
- max\_depth = 3
- subsample = 0.8
- scale\_pos\_weight = 1

**Accuracy = 79.11%**

**Precision score 0.72**

**Recall score 0.51**

**F1 score 0.59**

**AUC = 0.84**

**Confusion matrix :**

	0	1
0	287 (True -ve)	27 (False +ve)
1	67 (False -ve)	69 (True +ve)

**The XGBoost model from scratch gives an accuracy of 78.67% .**

# Comparative analysis

	Adaboost	XGBoost																		
Accuracy score	79.56%	79.11%																		
Precision score	0.72	0.72																		
Recall score	0.54	0.51																		
F1 score	0.62	0.59																		
AUC	0.86	0.84																		
Confusion matrix	<table border="1"><tr><td></td><td>0</td><td>1</td></tr><tr><td>0</td><td>285</td><td>29</td></tr><tr><td>1</td><td>63</td><td>73</td></tr></table>		0	1	0	285	29	1	63	73	<table border="1"><tr><td></td><td>0</td><td>1</td></tr><tr><td>0</td><td>287</td><td>27</td></tr><tr><td>1</td><td>67</td><td>69</td></tr></table>		0	1	0	287	27	1	67	69
	0	1																		
0	285	29																		
1	63	73																		
	0	1																		
0	287	27																		
1	67	69																		

We can see that both the models perform in a similar manner, with **Adaboost being slightly better than XGBoost**.

- Both the models have similar overall accuracy and precision.
- AdaBoost outperforms XGBoost on Recall score , F1 score, and ROC-AUC, indicating better ability to capture true positives and more effective class separation (also visible from confusion matrix) .
- Higher AUC in AdaBoost indicates better discrimination between classes.

We are detecting customer purchase so false positives must be avoided as marketing to uninterested users is costly. High precision models are required.  
**Both the models have precision > 0.7** .

- XGBoost may not have performed well due to class imbalance.
- The slightly better performance of AdaBoost could be attributed to the relatively small size of the dataset.

- Another reason can be less overfitting in AdaBoost as it uses Decision stumps(shallow decision trees) whereas XGBoost uses full Decision trees
- XGBoost has a lot of hyperparameters so tuning is more difficult than tuning in AdaBoost.